

GPU Programming using Python

(PyCUDA and PyOpenCL)

Ki-Hwan Kim

Department of Physics, Korea University

2010/5/28

Outline

- ① Why Python for scientific computations?
 - Scientific computing environments
- ② Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- ③ Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

Outline

- 1 Why Python for scientific computations?
 - Scientific computing environments
- 2 Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- 3 Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

Scientific computing environments

Matlab, Mathematica, Maple, ...

- simple and clean syntax
- high productivity
- tight integration of calculation and visualization

=> Do not work for some problems, at least not in an easy way

Traditional languages (Fortran, C, C++, JAVA)

- flexible and versatile
- high performance

=> Complicated work, Low productivity

Scientific computing environments

Matlab, Mathematica, Maple, ...

- simple and clean syntax
- high productivity
- tight integration of calculation and visualization

=> Do not work for some problems, at least not in an easy way

Traditional languages (Fortran, C, C++, JAVA)

- flexible and versatile
- high performance

=> Complicated work, Low productivity

Scientific computing environments

Matlab, Mathematica, Maple, ...

- simple and clean syntax
- high productivity
- tight integration of calculation and visualization

=> Do not work for some problems, at least not in an easy way

Traditional languages (Fortran, C, C++, JAVA)

- flexible and versatile
- high performance

=> Complicated work, Low productivity

Scientific computing environments

Matlab, Mathematica, Maple, ...

- simple and clean syntax
- high productivity
- tight integration of calculation and visualization

=> Do not work for some problems, at least not in an easy way

Traditional languages (Fortran, C, C++, JAVA)

- flexible and versatile
- high performance

=> Complicated work, Low productivity

Scientific computing environments (Cont.)

Scripting using Python

- simple and clean syntax
- gluing other favorite simulation, visualization, data analysis
- high productivity
- high performance (gluing Fortran, C)
- Easy integration with GPU computing (PyCUDA, PyOpenCL)

=> Build my own Matlab-like scientific computing environment,
tailored to my specific needs with high performance

Scientific computing environments (Cont.)

Scripting using Python

- simple and clean syntax
- gluing other favorite simulation, visualization, data analysis
- high productivity
- high performance (gluing Fortran, C)
- Easy integration with GPU computing (PyCUDA, PyOpenCL)

=> Build my own Matlab-like scientific computing environment,
tailored to my specific needs with high performance

Outline

- 1 Why Python for scientific computations?
 - Scientific computing environments
- 2 Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- 3 Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

Mathematical Formula

Gaussian and Sin functions

$$\begin{cases} a(x) &= e^{-\frac{x^2}{2}} \\ b(x) &= \sin(5x) \end{cases}$$

$$\Rightarrow c(x) = a(x) \times b(x) = e^{-\frac{x^2}{2}} \sin(5x)$$

Discretize Condition

$$-5 \leq x < 5 \quad \text{and} \quad \Delta_x = 0.01$$

$$\rightarrow n = 1000$$

Mathematical Formula

Gaussian and Sin functions

$$\begin{cases} a(x) &= e^{-\frac{x^2}{2}} \\ b(x) &= \sin(5x) \end{cases}$$

$$\Rightarrow c(x) = a(x) \times b(x) = e^{-\frac{x^2}{2}} \sin(5x)$$

Discretize Condition

$$-5 \leq x < 5 \quad \text{and} \quad \Delta_x = 0.01$$

$$\rightarrow n = 1000$$

Outline

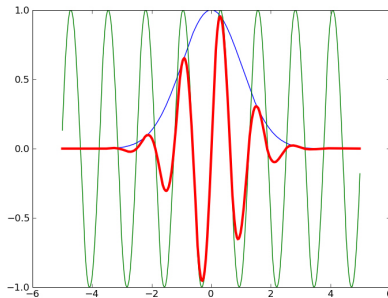
- 1 Why Python for scientific computations?
 - Scientific computing environments
- 2 Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- 3 Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

Python

```
1 #!/usr/bin/env python
2 import numpy as np
3
4 x = np.arange(-5,5,0.01,'f')
5 a = np.exp(-(x**2)/2)
6 b = np.sin(5*x)
7
8 c = a*b
```

Python (Plot the graph)

```
1 import matplotlib.pyplot as plt
2
3 plt.plot(x,a,x,b)
4 plt.plot(x,c,linewidth=3)
5 plt.savefig('./pics/mul_array.png')
6 plt.show()
```



PyCUDA

```
1 #!/usr/bin/env python
2 import numpy as np
3
4 x = np.arange(-5,5,0.01, 'f')
5 a = np.exp(-(x**2)/2)
6 b = np.sin(5*x)
7
8 import pycuda.driver as cuda
9 import pycuda.autoinit
10
11 a_gpu = cuda.to_device(a)
12 b_gpu = cuda.to_device(b)
13 c_gpu = cuda.mem_alloc(a.nbytes)
14
15 gpu_mul(np.int32(a.size), a_gpu, b_gpu, c_gpu,
16         block=(256,1,1), grid=(4,1))
17
18 c = np.zeros_like(a)
19 cuda.memcpy_dtoh(c, c_gpu)
```


PyCUDA (Cont.)

```
1 kernels = """
2     __global__ void multiply(int nx, float *a,
3         float *b, float *c){
4
5         int idx = threadIdx.x;
6
7         if(idx<nx) c[idx] = a[idx]*b[idx];
8     }
9 """
10
11 from pycuda.compiler import SourceModule
12 mod = SourceModule(kernels)
13 gpu_mul = mod.get_function("multiply")
```

PyOpenCL

```
1  #!/usr/bin/env python
2  import numpy as np
3
4  x = np.arange(-5,5,0.01, 'f')
5  a = np.exp(-(x**2)/2)
6  b = np.sin(5*x)
7
8  import pyopencl as cl
9  ctx = cl.create_some_context()
10 queue = cl.CommandQueue(ctx)
11
12 mf=cl.mem_flags
13 a_gpu = cl.Buffer(ctx, mf.COPY_HOST_PTR, hostbuf=a)
14 b_gpu = cl.Buffer(ctx, mf.COPY_HOST_PTR, hostbuf=b)
15 c_gpu = cl.Buffer(ctx, mf.COPY_HOST_PTR, size=a.nbytes)
16
17 gpu_mul(queue, (a.size,), np.int32(a.size), a_gpu, b_gpu, c_gpu)
18
19 c = np.zeros_like(a)
20 cl.enqueue_read_buffer(queue, c_gpu, c)
```

PyOpenCL (Cont.)

```
1 kernels = """
2     __kernel void multiply(int nx, __global float *a,
3         __global float *b, __global float *c) {
4
5         int idx = get_global_id(0);
6
7         if(idx<nx) c[idx] = a[idx]*b[idx];
8     }
9 """
10
11 prg = cl.Program(ctx, kernels).build()
12 gpu_mul = prg.multiply
```

Outline

- 1 Why Python for scientific computations?
 - Scientific computing environments
- 2 Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- 3 Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

PyCUDA vs PyOpenCL

```
1 import pycuda.driver as cuda
2 import pycuda.autotinit
3
4 a_gpu = cuda.to_device(a)
5 b_gpu = cuda.to_device(b)
6 c_gpu = cuda.mem_alloc(a.nbytes)
7 gpu_mul(np.int32(a.size), a_gpu, b_gpu, c_gpu,
8         block=(256,1,1), grid=(4,1))
9 cuda.memcpy_dtoh(c, c_gpu)
```

```
1 import pyopencl as cl
2 ctx = cl.create_some_context()
3 queue = cl.CommandQueue(ctx)
4
5 mf=cl.mem_flags
6 a_gpu = cl.Buffer(ctx, mf.COPY_HOST_PTR, hostbuf=a)
7 b_gpu = cl.Buffer(ctx, mf.COPY_HOST_PTR, hostbuf=b)
8 c_gpu = cl.Buffer(ctx, mf.COPY_HOST_PTR, size=a.nbytes)
9 gpu_mul(queue, (a.size,) np.int32(a.size), a_gpu, b_gpu, c_gpu)
10 cl.enqueue_read_buffer(queue, c_gpu, c)
```

PyCUDA vs PyOpenCL (Cont.)

```

1 kernels = """
2     __global__ void multiply(int nx, float *a,...) {
3         int idx = threadIdx.x;
4         if(idx<nx) c[idx] = a[idx]*b[idx];
5     }
6 """
7
8 from pycuda.compiler import SourceModule
9 mod = SourceModule(kernels)
10 gpu_mul = mod.get_function("multiply")

```

```

1 kernels = """
2     __kernel void multiply(int nx, __global float *a,...) {
3         int idx = get_global_id(0);
4         if(idx<nx) c[idx] = a[idx]*b[idx];
5     }
6 """
7
8 prg = cl.Program(ctx, kernels).build()
9 gpu_mul = prg.multiply

```

Outline

- ① Why Python for scientific computations?
 - Scientific computing environments
- ② Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- ③ Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

Mathematical Formula

2-D Wave Equation

$$\frac{\partial^2 f}{\partial t^2} = c^2 \nabla^2 f$$
$$\rightarrow \frac{\partial^2 f}{\partial t^2} = c^2 \left(\frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} \right)$$

Simple Finite-Difference Scheme

$$\frac{f_{i,j}^{n+1} - 2f_{i,j}^n + f_{i,j}^{n-1}}{\Delta_t^2} = c_{i,j}^2 \left(\frac{f_{i+1,j}^n - 2f_{i,j}^n + f_{i-1,j}^n}{\Delta_x^2} + \frac{f_{i,j+1}^n - 2f_{i,j}^n + f_{i,j-1}^n}{\Delta_y^2} \right)$$

$$f_{i,j}^{n+1} = \tilde{c}_{i,j}^2 (f_{i+1,j}^n + f_{i-1,j}^n + f_{i,j+1}^n + f_{i,j-1}^n - 4f_{i,j}^n) + 2f_{i,j}^n - f_{i,j}^{n-1}$$

Mathematical Formula

2-D Wave Equation

$$\frac{\partial^2 f}{\partial t^2} = c^2 \nabla^2 f$$
$$\rightarrow \frac{\partial^2 f}{\partial t^2} = c^2 \left(\frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} \right)$$

Simple Finite-Difference Scheme

$$\frac{f_{i,j}^{n+1} - 2f_{i,j}^n + f_{i,j}^{n-1}}{\Delta_t^2} = c_{i,j}^2 \left(\frac{f_{i+1,j}^n - 2f_{i,j}^n + f_{i-1,j}^n}{\Delta_x^2} + \frac{f_{i,j+1}^n - 2f_{i,j}^n + f_{i,j-1}^n}{\Delta_y^2} \right)$$

$$f_{i,j}^{n+1} = \tilde{c}_{i,j}^2 (f_{i+1,j}^n + f_{i-1,j}^n + f_{i,j+1}^n + f_{i,j-1}^n - 4f_{i,j}^n) + 2f_{i,j}^n - f_{i,j}^{n-1}$$

Outline

- ① Why Python for scientific computations?
 - Scientific computing environments
- ② Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- ③ Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

Python

```
1 import numpy as np
2
3 c = np.ones((1000,1000), 'f')*0.25
4 f = np.zeros_like(c)
5 g = np.zeros_like(c)
6
7 ii = (slice(1,-1), slice(1,-1))
8 for tn in xrange(1000):
9     g[400,500] += np.sin(0.1*tn)
10    f[ii] = c[ii]*(g[2:,1:-1]+g[:-2,1:-1]+g[1:-1,2:]+g[1:-1,:-2]
11               -4g[ii])+2g[ii]-f[ii]
12    g[ii] = c[ii]*(f[2:,1:-1]+f[:-2,1:-1]+f[1:-1,2:]+f[1:-1,:-2]
13               -4f[ii])+2f[ii]-g[ii]
```

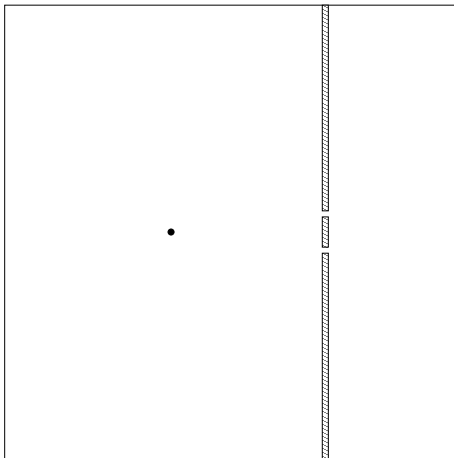
PyCUDA

```
1 kernels = """
2     __global__ void update_src(int idx, int tn, float *g) {
3         g[idx] += sin(0.1*tn);
4     }
5
6     __global__ void update(int nx, int ny, float *c, float *f,
7         float *g) {
8         int idx = blockIdx.x*blockDim.x + threadIdx.x;
9         int i = idx/ny, j = idx%ny;
10
11         if(i>0 && j>0 && i<nx-1 && j<ny-1)
12             f[idx] = c[idx]*(g[idx+ny]+g[idx-ny]+g[idx+1]+g[idx-1]
13                 -4*g[idx])+2*g[idx]-f[idx];
14     }
15 """
16 from pycuda.compiler import SourceModule
17 mod = SourceModule(kernels)
18 update_src = mod.get_function("update_src")
19 update = mod.get_function("update")
```

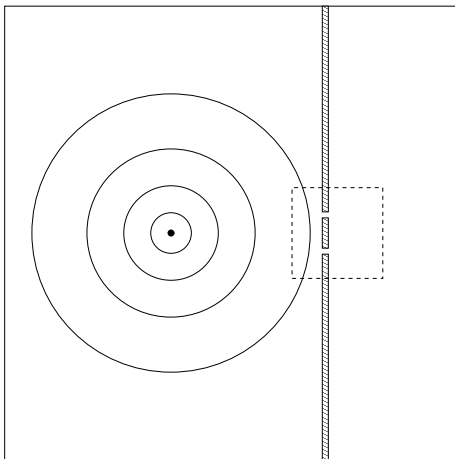
PyCUDA (Cont.)

```
1 import pycuda.driver as cuda
2 import pycuda.autoinit
3
4 c = np.ones((nx,ny), 'f')*0.25
5 f = np.zeros_like(c)
6 c_gpu = cuda.to_device(c)
7 f_gpu = cuda.to_device(f)
8 g_gpu = cuda.to_device(f)
9 ...
10 cuda.memcpy_htod(c_gpu, c)
11
12 Db, Dg = (256,1,1), (nx*ny/256+1,1)
13 nnx, nny = np.int32(nx), np.int32(ny)
14 src_pt = np.int32(nx/2*ny + ny/2)
15 ...
16 for tn in xrange(10000):
17     update_src(src_pt, np.int32(tn), f_gpu, (1,1,1), (1,1))
18     update(nnx, nny, c_gpu, f_gpu, g_gpu, Db, Dg)
19     update(nnx, nny, c_gpu, g_gpu, f_gpu, Db, Dg)
20     if tn%100==0: cuda.memcpy_dtoh(f, f_gpu)
```

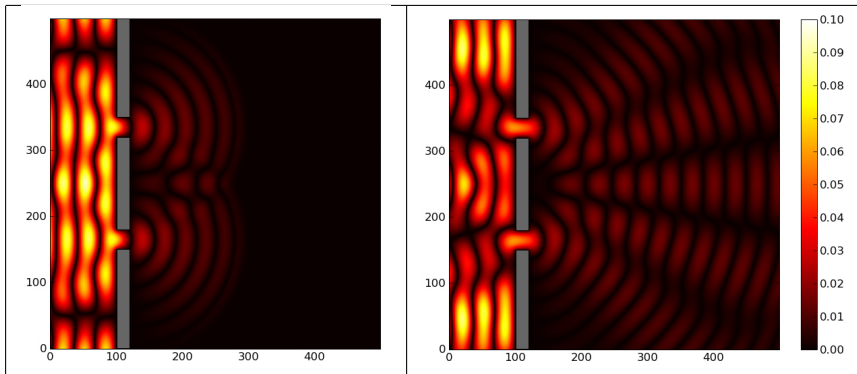
Example - Double Slit Interference



Double Slit Interference (source point)



Double Slit Interference (result)



Outline

- ① Why Python for scientific computations?
 - Scientific computing environments
- ② Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- ③ Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - **Shared Memory Optimization**
 - Utilize Multi-GPUs with MPI

Shared Memory Optimize

```
1  __global__ void update(int nx, int ny, float *c, float *f,  
    float *g) {  
2      int idx = blockIdx.x*blockDim.x + threadIdx.x;  
3      int i = idx/ny, j = idx%ny;  
4  
5      if(i>0 && j>0 && i<nx-1 && j<ny-1)  
6          f[idx] = c[idx]*(g[idx+ny]+g[idx-ny]  
7              +g[idx+1]+g[idx-1]-4*g[idx])+2*g[idx]-f[idx];  
8  }
```

Duplicated arrays in a thread block

=> transfer to **shared memory**, and reuse

=> reduce **global memory access**

Target: g[idx+1], g[idx-1], g[idx], g[idx]

Shared Memory Optimize (Cont.)

```
1  __global__ void update(int nx, int ny, float *c, float *f,  
    float *g) {  
2      int tx = threadIdx.x;  
3      int idx = blockIdx.x*blockDim.x + tx;  
4  
5      extern __shared__ float gs[];  
6      gs[tx+1] = g[idx];  
7      int i = idx/ny, j = idx%ny;  
8      if(j>0 && j<ny-1) {  
9          if(tx==0) gs[tx] = g[idx-1];  
10         if(tx==blockDim.x-1) gs[tx+2] = g[idx+1];  
11     }  
12     __syncthreads()  
13  
14     if(i>0 && j>0 && i<nx-1 && j<ny-1)  
15         f[idx] = c[idx]*(g[idx+ny]+g[idx-ny]  
16             +gs[tx+2]+gs[tx]-4*gs[tx+1])+2*gs[tx+1]-f[idx];  
17 }
```

```
1  update(nnx, nny, c_gpu, f_gpu, g_gpu, Db, Dg, shared=(256+2)*4)  
2  update(nnx, nny, c_gpu, g_gpu, f_gpu, Db, Dg, shared=(256+2)*4)
```

Shared Memory Optimize (Cont.)

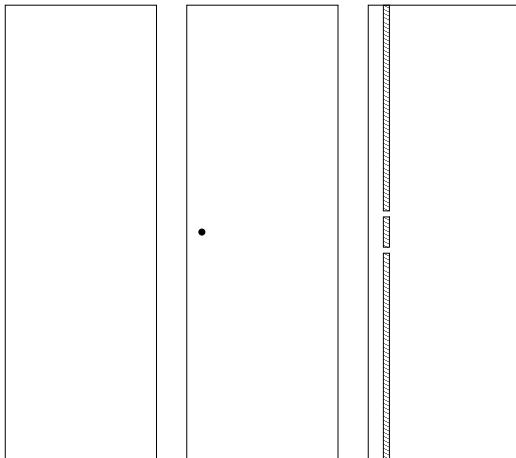
```
1  __global__ void update(int nx, int ny, float *c, float *f,  
    float *g) {  
2      int tx = threadIdx.x;  
3      int idx = blockIdx.x*blockDim.x + tx;  
4  
5      extern __shared__ float gs[];  
6      gs[tx+1] = g[idx];  
7      int i = idx/ny, j = idx%ny;  
8      if(j>0 && j<ny-1) {  
9          if(tx==0) gs[tx] = g[idx-1];  
10         if(tx==blockDim.x-1) gs[tx+2] = g[idx+1];  
11     }  
12     __syncthreads()  
13  
14     if(i>0 && j>0 && i<nx-1 && j<ny-1)  
15         f[idx] = c[idx]*(g[idx+ny]+g[idx-ny]  
16             +gs[tx+2]+gs[tx]-4*gs[tx+1])+2*gs[tx+1]-f[idx];  
17 }
```

```
1  update(nnx, nny, c_gpu, f_gpu, g_gpu, Db, Dg, shared=(256+2)*4)  
2  update(nnx, nny, c_gpu, g_gpu, f_gpu, Db, Dg, shared=(256+2)*4)
```

Outline

- ① Why Python for scientific computations?
 - Scientific computing environments
- ② Example 1) 1-D Array Multiplication
 - Formula
 - Python, PyCUDA, PyOpenCL
 - PyCUDA vs PyOpenCL
- ③ Example 2) 2-D Wave Equation
 - Formula
 - Python, PyCUDA
 - Shared Memory Optimization
 - Utilize Multi-GPUs with MPI

Divide Domain



Move to MPI

```
import pycuda.driver as cuda
import pycuda.autotinit

c = sc.ones((nx,ny), 'f')*0.25
f = sc.zeros_like(c)
c_gpu = cuda.to_device(c)
f_gpu = cuda.to_device(f)
g_gpu = cuda.to_device(f)

# set the c array with geometry
cuda.memcpy_htod(c_gpu, c)

Db, Dg = (256,1,1), (nx*ny/256+1,1)
nnx, nny = sc.int32(nx), sc.int32(ny)
src_pt = sc.int32(nx/2*ny + ny/2)

# SourceModule

for tn in xrange(10000):
    update_src(src_pt, ...)
    update(..., f_gpu, g_gpu, ...)
    update(..., g_gpu, f_gpu, ...)

    if tn%100==0:
        cuda.memcpy_dtoh(f, f_gpu)
```

```
import pycuda.driver as cuda
import boostmpi as mpi

cuda.init()
dev = cuda.Device(mpi.rank)
ctx = dev.make_context()

# memory allocate

if mpi.rank == 2:
    # set the c array with geometry
    cuda.memcpy_htod(c_gpu, c)

Db, Dg = ...
nnx, nny = ...
if mpi.rank == 1:
    src_pt = ...

# SourceModule

for tn in xrange(10000):
    if mpi.rank == 1:
        update_src(src_pt, ...)
        update(..., f_gpu, g_gpu, ...)
        update(..., g_gpu, f_gpu, ...)

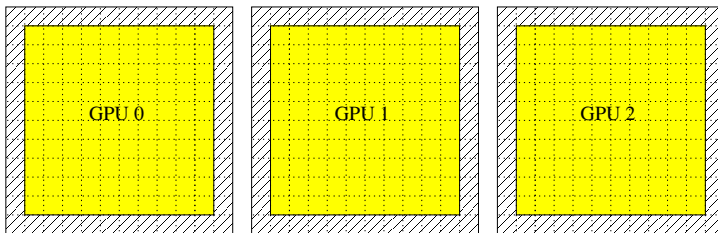
ctx.pop()
```

Gather data

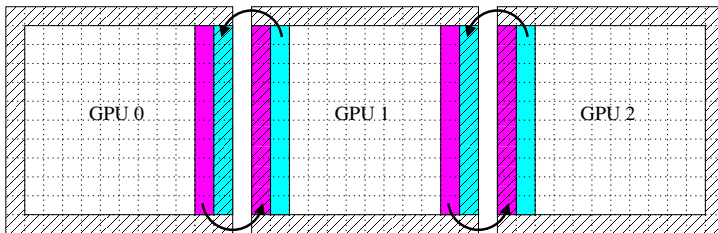
```
1  nxx, nyy = nx-2, ny-2
2
3  if mpi.rank == 0:
4      output = sc.zeros(3*nxx, nyy, 'f')
5
6  for tn in xrange(1000):
7      ...
8      if tn > 1000:
9          if mpi.rank == 0:
10             cuda.memcpy_dtoh(f, f_gpu)
11             output[:nxx,:] = f[1:-1,1:-1]
12             output[nxx:2*nxx,:] = mpi.world.recv(1,3)
13             output[2*nxx:3*nxx,:] = mpi.world.recv(2,3)
14         else:
15             cuda.memcpy_dtoh(f, f_gpu)
16             mpi.world.send(0, 3, f[1:-1,1:-1])
```


Exchange Boundaries

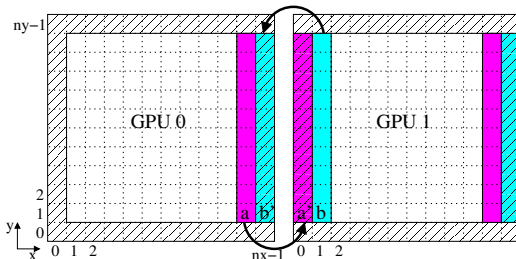
STEP 1



STEP 2



Memcpy Offset



offset

$$a: \text{int}(a_gpu) + ((nx-2)*ny+1)*sof$$

$$a': \text{int}(a_gpu) + 1*sof$$

$$b: \text{int}(a_gpu) + (ny+1)*sof$$

$$b': \text{int}(a_gpu) + ((nx-1)*ny+1)*sof$$

Exchange Boundaries

```
1 nbof = np.nbytes[ 'float32 ' ]
2 dtof = np.dtype( 'float32 ' )
3
4 def send(rank, tag_mark, nx, ny, a_gpu):
5     if mpi.rank < rank:
6         offset = int(a_gpu)+((nx-2)*ny+1)*nbof
7     elif mpi.rank > rank:
8         offset = int(a_gpu)+(ny+1)*nbof
9     mpi.world.send(rank, tag_mark, \
10                    cuda.from_device(offset, (ny-2,), dtof) )
11
12 def recv(rank, tag_mark, nx, ny, a_gpu):
13     if mpi.rank > rank:
14         offset = int(a_gpu)+1*nbof
15     elif mpi.rank < rank:
16         offset = int(a_gpu)+((nx-1)*ny+1)*nbof
17     cuda.memcpy_htod(offset, mpi.world.recv(rank, tag_mark))
```

Exchange Boundaries (Cont.)

```
1 def exchange(nx, ny, a_gpu):
2     if mpi.rank == 0:
3         send(1, 0, nx, ny, a_gpu)
4         recv(1, 0, nx, ny, a_gpu)
5     if mpi.rank == 1:
6         recv(0, 0, nx, ny, a_gpu)
7         send(0, 1, nx, ny, a_gpu)
8         send(2, 0, nx, ny, a_gpu)
9         recv(2, 1, nx, ny, a_gpu)
10    if mpi.rank == 2:
11        recv(1, 0, nx, ny, a_gpu)
12        send(1, 1, nx, ny, a_gpu)
```

```
1 for tn in xrange(10000):
2     if mpi.rank == 1:
3         update_src(src_pt, ...)
4         update(..., f_gpu, g_gpu, ...)
5         exchange(nx, ny, f_gpu)
6         update(..., g_gpu, f_gpu, ...)
7         exchange(nx, ny, g_gpu)
```

Have your Fun with Python!

THANK YOU