

# Scientific Computing with Python and CUDA

Stefan Reiterer

High Performance Computing Seminar, January 17 2011

# Inhalt

- 1 A short Introduction to Python
- 2 Scientific Computing tools in Python
- 3 Easy ways to make Python faster
- 4 Python + CUDA = PyCUDA
- 5 Python + MPI = mpi4py
- 6 ...and the Rest

# Inhalt

- 1 A short Introduction to Python
- 2 Scientific Computing tools in Python
- 3 Easy ways to make Python faster
- 4 Python + CUDA = PyCUDA
- 5 Python + MPI = mpi4py
- 6 ...and the Rest

# What is Python?

Python is

# What is Python?

Python is

- a **high level** programming language

# What is Python?

Python is

- a **high level** programming language
- **interpreted**,

# What is Python?

Python is

- a **high level** programming language
- **interpreted**,
- **object oriented**,

# What is Python?

Python is

- a **high level** programming language
- **interpreted**,
- **object oriented**,
- named after the "Monty Pythons".



# What is Python?

Python is

- a **high level** programming language
- **interpreted**,
- **object oriented**,
- named after the "Monty Pythons".

It was invented by Guido van Rossum in the early 90s.

# Why Python?

- Object Oriented programming

# Why Python?

- Object Oriented programming → reusable code

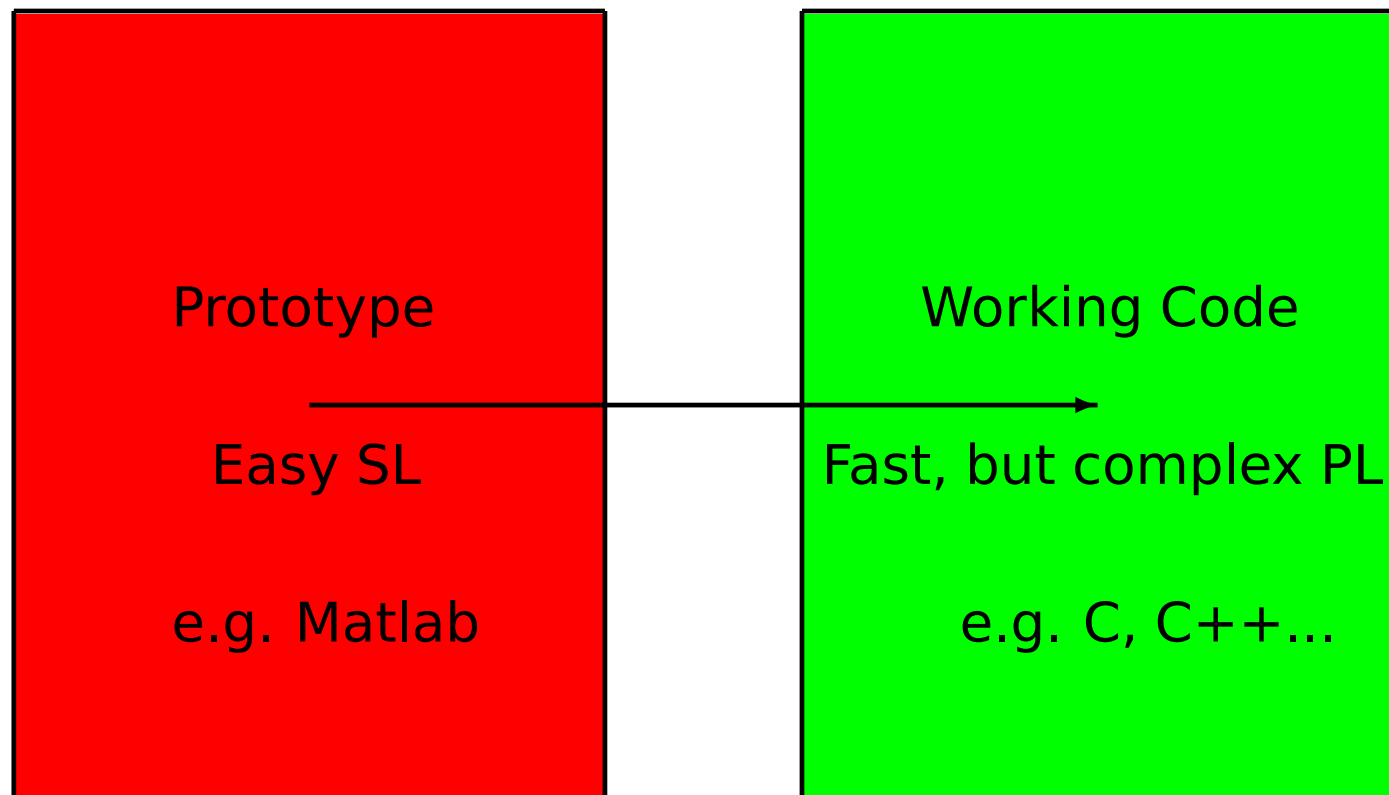
# Why Python?

- Object Oriented programming → reusable code
- True strength of Python: Rapid prototyping

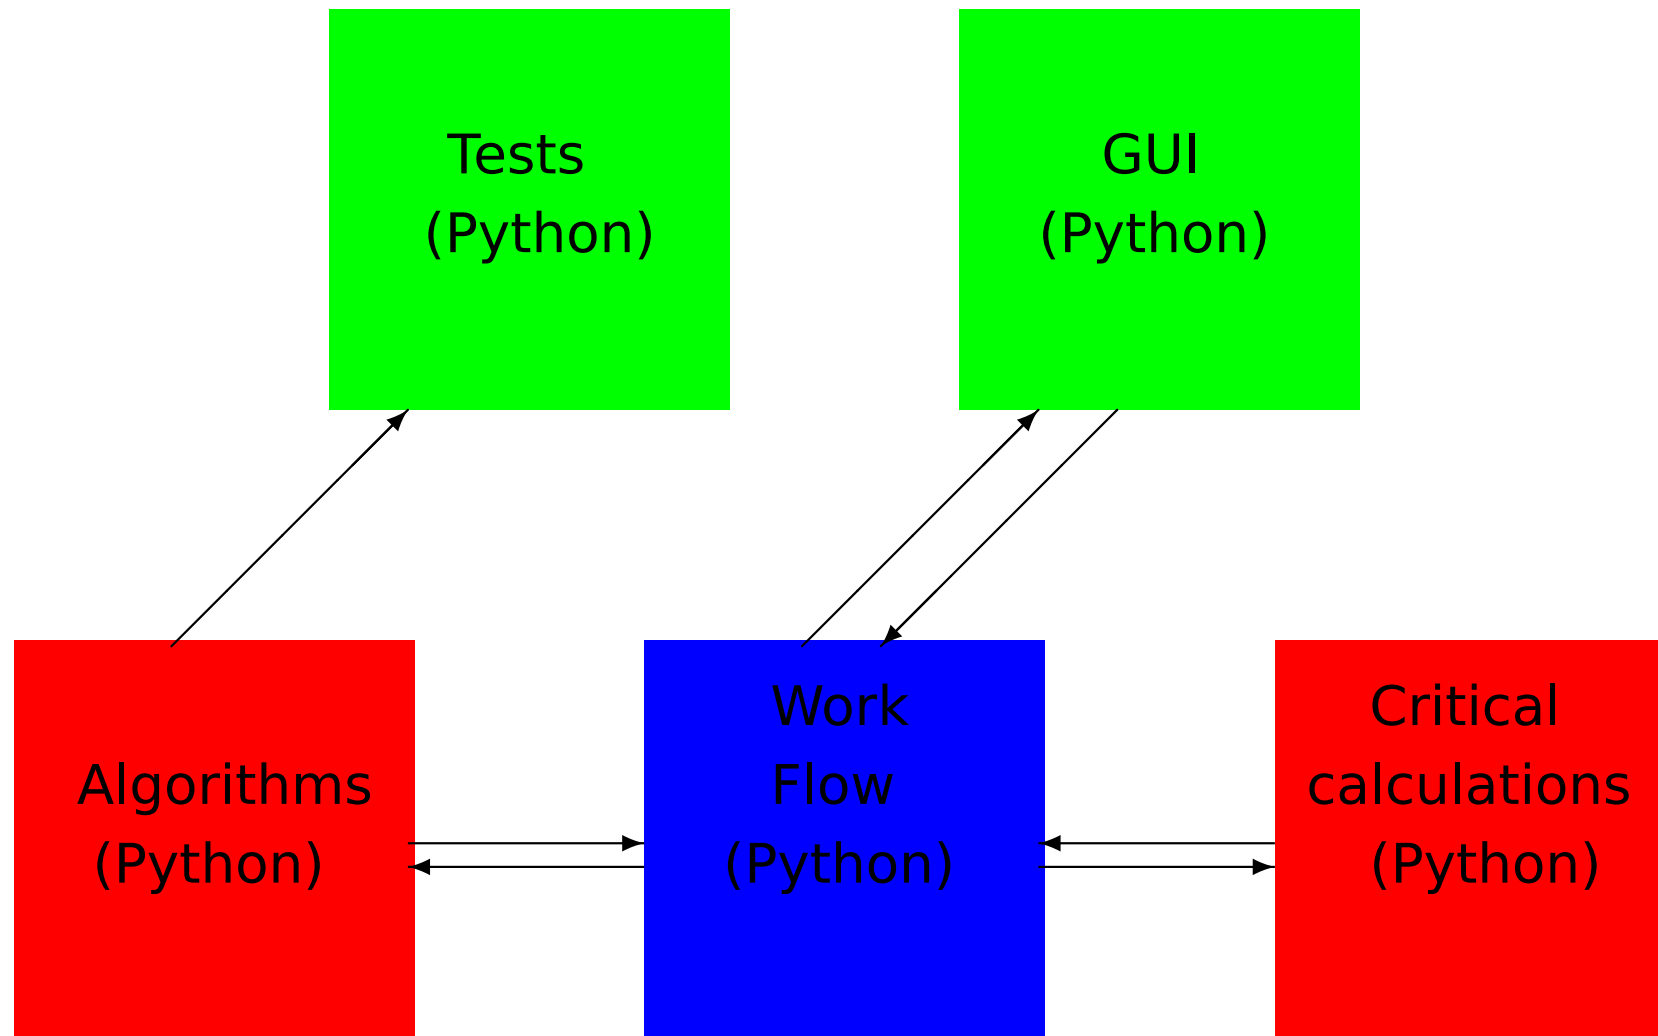
# Why Python?

- Object Oriented programming → reusable code
- True strength of Python: Rapid prototyping → accelerate research process, saves time and nerves.

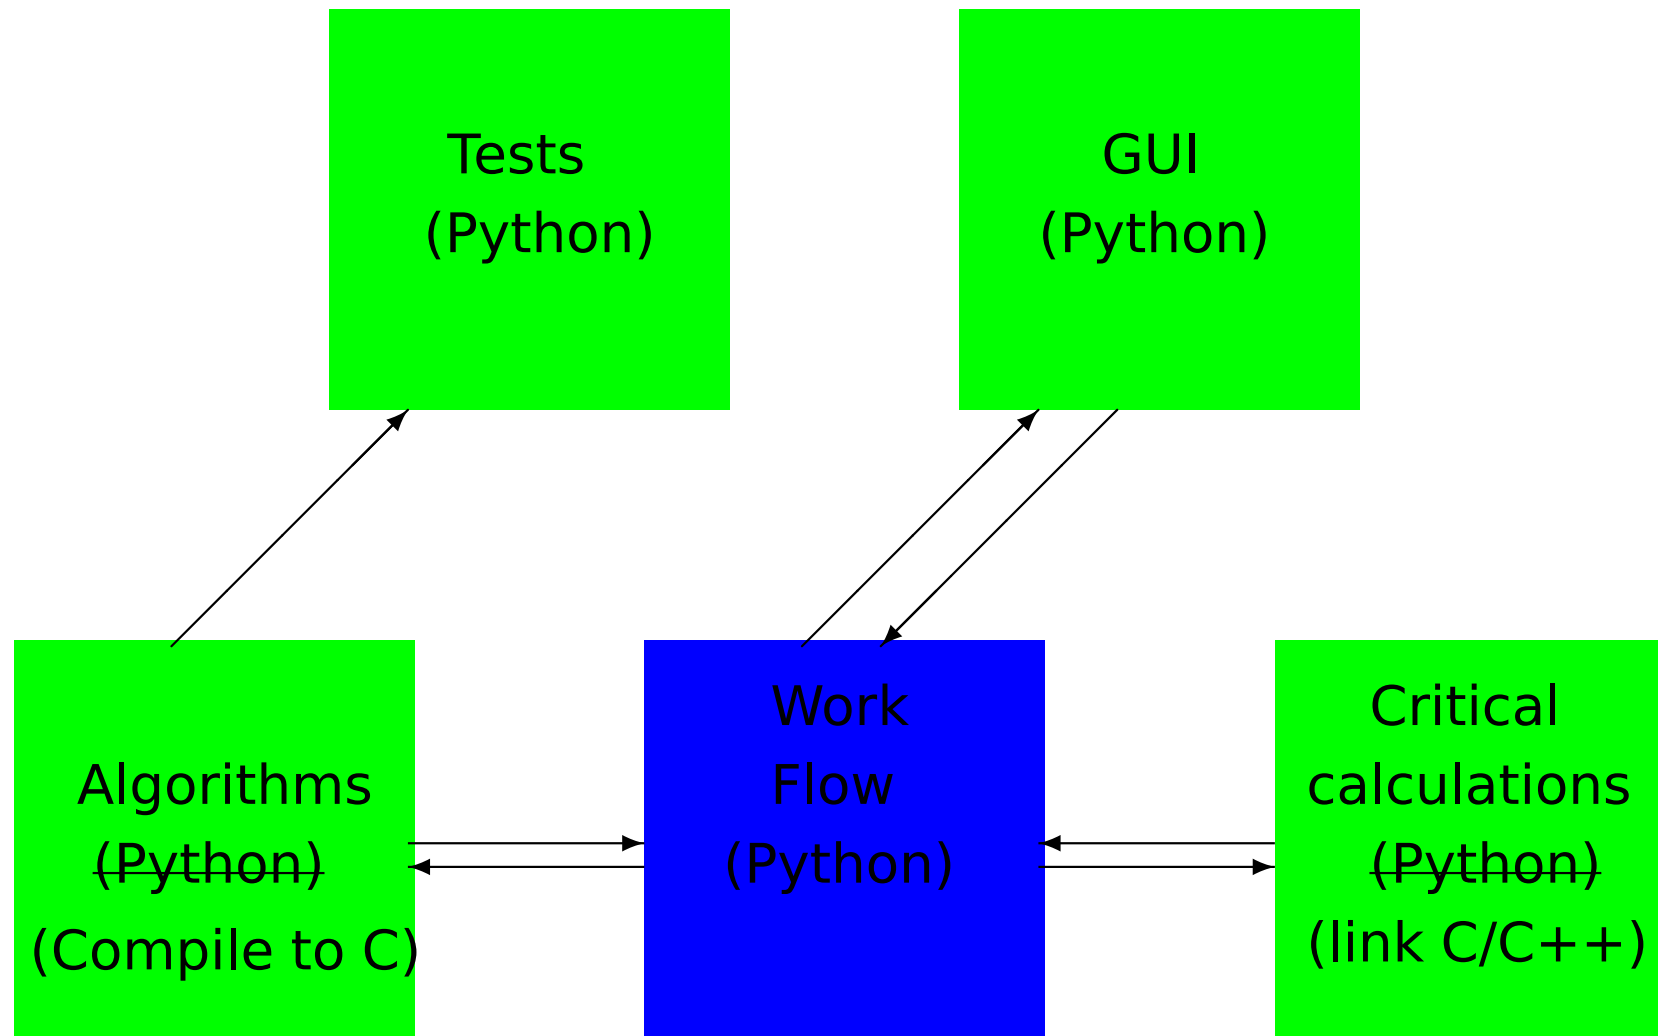
# Classical prototyping process



# Prototyping with Python



# Prototyping with Python





# Don't fear Python!



# Don't fear Python!

- Intuitive Syntax
- No Allocation
- Garbage Collection
- Easy to Learn!



# The Goodbye World Program

```
print( "GoodbyeWorld! " )
```

# Definition of functions

```
def function_name(arg1,arg2,arg3 = default):  
    result = arg1 + arg2*arg3 #do something  
    return result
```

# Definition of functions

```
def function_name(arg1, arg2, arg3 = default):  
    result = arg1 + arg2*arg3 #do something  
    return result
```

Python recognizes code blocks with **intends**!

# Definition of functions

```
def function_name(arg1, arg2, arg3 = default):  
    result = arg1 + arg2*arg3 #do something  
    return result
```

Python recognizes code blocks with **intends**!

**Always remember:** Don't mix tabs with spaces!

# If statements

```
if cond1 is True: #normal if
    do this
elif cond2 is True: #else if
    do that
else: #else
    do something_else
```

# while loops

```
while condition is True:  
    do something
```



# while loops

```
while condition is True:  
    do something
```

While loops know **else** too:

```
while condition is True:  
    do something  
else:  
    do something_else
```

# for loops

```
for x in range(n):  
    do something with x
```

# for loops

```
for x in range(n):  
    do something with x
```

- Instead of **range** one can use **any list**!

# for loops

```
for x in range(n):  
    do something with x
```

- Instead of **range** one can use **any list**!
- For better performance use **xrange**.

# for loops

```
for x in range(n):  
    do something with x
```

- Instead of **range** one can use **any list**!
- For better performance use **xrange**.
- for loops also have an **else** statement.

## An example...

```
for n in range(2, 10):  
    for x in range(2, n):  
        if n % x == 0:  
            print n, 'equals', x, '*', n/x  
            break  
    else:  
        # loop fell through without finding a factor  
        print n, 'is_a_prime_number'
```

# Output

```
2 is a prime number
3 is a prime number
4 equals 2 * 2
5 is a prime number
6 equals 2 * 3
7 is a prime number
8 equals 2 * 4
9 equals 3 * 3
```

# Classes

In Python **everything** is a class!



# Classes

In Python **everything** is a class!  
**No** elementary data types!

# Definition of a class

**Classes** are defined in the same way as functions

```
class my_class:  
    statement1  
    statement2  
    ...
```

# Definition of a class

**Classes** are defined in the same way as functions

```
class my_class:  
    statement1  
    statement2  
    ...
```

- **`__init__`** serves as **constructor**

# Definition of a class

**Classes** are defined in the same way as functions

```
class my_class:  
    statement1  
    statement2  
    ...
```

- **`__init__`** serves as **constructor**
- **`__del__`** serves as **destructor**

# Definition of a class

**Classes** are defined in the same way as functions

```
class my_class:  
    statement1  
    statement2  
    ...
```

- **`__init__`** serves as **constructor**
- **`__del__`** serves as **destructor**
- Class methods take the object itself as first argument (as convention one writes **`self`**)

# Derivative of classes

Classes are derived simply with

```
class my_class(inheritance):  
    statement1  
    statement2  
    ...
```

# Derivative of classes

Classes are derived simply with

```
class my_class(inheritance):  
    statement1  
    statement2  
    ...
```

- They hold the same methods and variables as their parents

# Derivative of classes

Classes are derived simply with

```
class my_class(inheritance):  
    statement1  
    statement2  
    ...
```

- They hold the same methods and variables as their parents
- To overload methods simply add the **same method again**.



# Operator overloading

To overload operators like `+`, `-`, `*`, `/` etc simply add the methods **`__add__`**, **`__sub__`**, **`__mul__`**, **`__div__`** ...

## Another example:

```
class band_matrix:
    def __init__(self, ab):
        self.shape = (ab.shape[1], ab.shape[1])
        self.data = ab
        self.band_width = ab.shape[0]
        self.dtype = ab.dtype

    def matvec(self, u):
        if self.shape[0] != u.shape[0]:
            raise ValueError("Dimension_Mismatch!")
        return band_matvec(self.data, u)

    def __add__(self, other):
        return add_band_mat(self, other)
```

# Inhalt

- 1 A short Introduction to Python
- 2 Scientific Computing tools in Python**
- 3 Easy ways to make Python faster
- 4 Python + CUDA = PyCUDA
- 5 Python + MPI = mpi4py
- 6 ...and the Rest

# NumPy and SciPy

- NumPy is for Numerical Python
- SciPy is for Scientific Python

# NumPy and SciPy

- NumPy is for Numerical Python
- SciPy is for Scientific Python
- Allow Numerical Linear Algebra like in Matlab

# NumPy and SciPy

- NumPy is for Numerical Python
- SciPy is for Scientific Python
- Allow Numerical Linear Algebra like in Matlab
- Speed: NumPy + SciPy  $\approx$  Matlab

# Some examples...

# Inhalt

- 1 A short Introduction to Python
- 2 Scientific Computing tools in Python
- 3 Easy ways to make Python faster**
- 4 Python + CUDA = PyCUDA
- 5 Python + MPI = mpi4py
- 6 ...and the Rest



# Band Matrix vector Multiplication

```
# -*- coding: utf-8 -*-  
def band_matvec(A,u):  
  
    result = zeros(u.shape[0],dtype=u.dtype)  
  
    for i in xrange(A.shape[1]):  
        result[i] = A[0,i]*u[i]  
  
    for j in xrange(1,A.shape[0]):  
        for i in xrange(A.shape[1]-j):  
            result[i] += A[j,i]*u[i+j]  
            result[i+j] += A[j,i]*u[i]  
  
    return result
```



# Write the C-Code inline in Python with Blitz

```
# -*- coding: utf-8 -*-  
from numpy import array, zeros  
from scipy.weave import converters  
from scipy import weave  
  
def band_matvec_inline(A,u):  
  
    result = zeros(u.shape[0],dtype=u.dtype)  
  
    N = A.shape[1]  
    B = A.shape[0]
```

# Write the C-Code inline in Python with Blitz

```
code = """
for(int i=0; i < N;i++) {
    result(i) = A(0,i)*u(i);
}
for(int j=1;j < B;j++) {
    for(int i=0; i < (N-j);i++) {
        if((i+j < N)) {
            result(i) += A(j,i)*u(j+i);
            result(i+j) += A(j,i)*u(i);
        }
    }
}
"""
```

# Write the C-Code inline in Python with Blitz

```
weave.inline(code,['u', 'A', 'result', 'N', 'B'],  
             type_converters=converters.blitz)  
return result
```

# Write the C-Code inline in Python with Blitz

```
weave.inline(code,['u', 'A', 'result', 'N', 'B'],  
             type_converters=converters.blitz)  
return result
```

≈ 290x speedup to python

# Cython

Python produces a lot of **overhead** with data types

# Cython

Python produces a lot of **overhead** with data types

**Solution:** Declare which data types to use and compile it with **Cython!**

# Python again

```
# -*- coding: utf-8 -*-  
def band_matvec(A,u):  
  
    result = zeros(u.shape[0],dtype=u.dtype)  
  
    for i in xrange(A.shape[1]):  
        result[i] = A[0,i]*u[i]  
  
    for j in xrange(1,A.shape[0]):  
        for i in xrange(A.shape[1]-j):  
            result[i] += A[j,i]*u[i+j]  
            result[i+j] += A[j,i]*u[i]  
  
    return result
```





# Get more with Cython

```
import numpy, scipy
cimport numpy as cnumpy

ctypedef cnumpy.float64_t reals
def matvec_help(cnumpy.ndarray[reals, ndim=2] A,
                cnumpy.ndarray[reals, ndim=1] u):
    cdef Py_ssize_t i, j
    cdef cnumpy.ndarray[reals, ndim=1] result = \
        numpy.zeros(A.shape[1], dtype=A.dtype)
    for i in xrange(A.shape[1]):
        result[i] = A[0, i]*u[i]
    for j in xrange(1, A.shape[0]):
        for i in xrange(A.shape[1]-j):
            result[i] = result[i] + A[j, i]*u[i+j]
            result[i+j] = result[i+j]+A[j, i]*u[i]
```



## ...and even more

```
import numpy, scipy
import numpy as cnumpy, cython
@cython.boundscheck(False)
ctypedef cnumpy.float64_t reals
def matvec_help(cnumpy.ndarray[reals, ndim=2] A,
                cnumpy.ndarray[reals, ndim=1] u):
    cdef Py_ssize_t i, j
    cdef cnumpy.ndarray[reals, ndim=1] result = \
        numpy.zeros(A.shape[1], dtype=A.dtype)
    for i in xrange(A.shape[1]):
        result[i] = A[0, i] * u[i]
    for j in xrange(1, A.shape[0]):
        for i in xrange(A.shape[1] - j):
            result[i] = result[i] + A[j, i] * u[i+j]
            result[i+j] = result[i+j] + A[j, i] * u[i]
```



# Speedup to python

# Speedup to python

- Cython  $\approx 220\times$  speedup

# Speedup to python

- Cython  $\approx 220\times$  speedup
- ...with boundscheck:  $\approx 440\times$  speedup

# Overview

Language	time	speedup
Python	2.96 s	1×
Matlab	90 ms	30×
Pure C	50 ms	60×
Cython	12 ms	220×
Inline C++	10 ms	290×
Cython w. boundscheck	6 ms	440×

# Overview

Language	time	speedup
Python	2.96 s	1×
Matlab	90 ms	30×
Pure C	50 ms	60×
Cython	12 ms	220×
Inline C++	10 ms	290×
Cython w. boundscheck	6 ms	440×

With linking libraries like Lapack even **more** is possible!

# Inhalt

- 1 A short Introduction to Python
- 2 Scientific Computing tools in Python
- 3 Easy ways to make Python faster
- 4 Python + CUDA = PyCUDA**
- 5 Python + MPI = mpi4py
- 6 ...and the Rest



# What is PyCUDA

- PyCUDA is a Wrapper in Python to CUDA
- Inline CUDA
- Garbage Collection
- A NumPy like vector class.
- Developed by Andreas Klöckner

# Step 1: Write your Code in CUDA

```
source = """
__global__ void doublify(float *a)
{
    int idx = threadIdx.x + threadIdx.y*4;
    a[idx] *= 2;
}
"""
```

## Step 2: Get it into Python

```
import pycuda.driver as cuda
import pycuda.autoinit
from pycuda.compiler import SourceModule

mod = SourceModule(source)
func = mod.get_function("doublify")
```

## Step3: Call it

```
import numpy
#create vector
x = numpy.random.randn(4,4)
x = x.astype(numpy.float32)
#copy it on card
x_gpu = cuda.mem_alloc(x.nbytes)
cuda.memcpy_htod(x_gpu, x)
#call function
func(x_gpu, block=(4,4,1))
#get data back
x_doubled = numpy.empty_like(x)
cuda.memcpy_dtoh(x_doubled, x_gpu)
```

# Benefits

- You can compile CUDA code on runtime

# Benefits

- You can compile CUDA code on runtime
- You can call it in Python

# Benefits

- You can compile CUDA code on runtime
- You can call it in Python

I was able to implement CG algorithm with the different variants  
**without rewriting code!**

# Overview (Band-matrix vector multiplication)

Language	time	speedup
Python	2.87 s	1×
Cython	9.87 ms	270×
Cython w. boundscheck	5.56 ms	500×
PyCUDA (on GTX 280)	585 $\mu$ s	5000×



# Overview (Band-matrix vector multiplication)

Language	time	speedup
Python	2.87 s	1×
Cython	9.87 ms	270×
Cython w. boundscheck	5.56 ms	500×
PyCUDA (on GTX 280)	585 $\mu$ s	5000×

# Inhalt

- 1 A short Introduction to Python
- 2 Scientific Computing tools in Python
- 3 Easy ways to make Python faster
- 4 Python + CUDA = PyCUDA
- 5 Python + MPI = mpi4py**
- 6 ...and the Rest

## Now something completely different... mpi4py

```
from mpi4py import MPI
comm = MPI.COMM_WORLD
print("hello_world")
print("my_rank_is :_%d"%comm.rank)
```

## Now something completely different... mpi4py

```
from mpi4py import MPI
comm = MPI.COMM_WORLD
print("hello_world")
print("my_rank_is :_%d"%comm.rank)
```

Simply call Python with MPI:

```
mpirun -n <# processes> python filename.py
```

# PyCUDA+mpi4py

```
from mpi4py import MPI
import pycuda.driver as cuda
from pycuda import gpuarray
from numpy import float32, array
from numpy.random import randn as rand
```

# PyCUDA+mpi4py

```
comm = MPI.COMM_WORLD
rank = comm.rank
root = 0
nr_gpus = 4
sendbuf = []
N = 2**20*nr_gpus
K = 1000
if rank == 0:
    x = rand(N).astype(float32)*10**16
    print "x:", x
    cuda.init() #init pycudadriver
    sendbuf = x.reshape(nr_gpus,N/nr_gpus)
```

# PyCUDA+mpi4py

```
if rank > nr_gpus-1:  
    raise ValueError("To_few_gpus!")  
current_dev = cuda.Device(rank) #device we are working on  
ctx = current_dev.make_context() #make a working context  
ctx.push() #let context make the lead  
#recieve data and port it to gpu:  
x_gpu_part = gpuarray.to_gpu(comm.scatter(sendbuf,root))
```

# PyCUDA+mpi4py

```
#do something ...  
for k in xrange(K):  
    x_gpu_part = 0.9*x_gpu_part
```



# PyCUDA+mpi4py

```
#get data back:  
x_part = (x_gpu_part).get()  
ctx.pop() #deactivate again  
ctx.detach() #delete it
```

# PyCUDA+mpi4py

```
recvbuf=comm.gather(x_part,root) #recieve data
if rank == 0:
    x_altered= array(recvbuf).reshape(N)
    print "altered_x:", x_altered
```

# Overview

N	K	NumPy	MPI+Numpy	MPI+PyCUDA
$2^{17}$	500	125 ms	4.6 ms	3 s
$2^{18}$	1000	636 ms	212 ms	3.7 s
$2^{19}$	2000	2.68 s	781 ms	3.84 s
$2^{20}$	4000	10.9 s	6.11 s	7.6 s
$2^{21}$	8000	39.5 s	24.2 s	11.6 s
$2^{22}$	16000	30 min	93.0 s	19.3 s
$2^{23}$	32000	?	382 s	24.1 s
$2^{24}$	64000	?	25 min	48.0 s

# Inhalt

- 1 A short Introduction to Python
- 2 Scientific Computing tools in Python
- 3 Easy ways to make Python faster
- 4 Python + CUDA = PyCUDA
- 5 Python + MPI = mpi4py
- 6 ...and the Rest**

# Where can I get Python?

Everything is **Open**source

**Links:**

- <http://www.python.org/>
- <http://www.scipy.org/>
- <http://www.sagemath.org/>
- <http://femhub.org/>
- <http://cython.org/>

# Where can I get Python?

Everything is **Open**source

**Links:**

- <http://www.python.org/>
- <http://www.scipy.org/>
- <http://www.sagemath.org/>
- <http://femhub.org/>
- <http://cython.org/>

If you need the latest mpi4py and and PyCUDA packages for Sage/Femhub ask me!

# Other interesting things for Python

- **f2py** → Fortran Code inline
- **petsc4py** → PETSC interface for Python
- **PyOpenCL** → like PyCUDA but for OpenCL
- **matplotlib, Mayavi** → powerful plotting libraries

... and many many more

# Questions?

